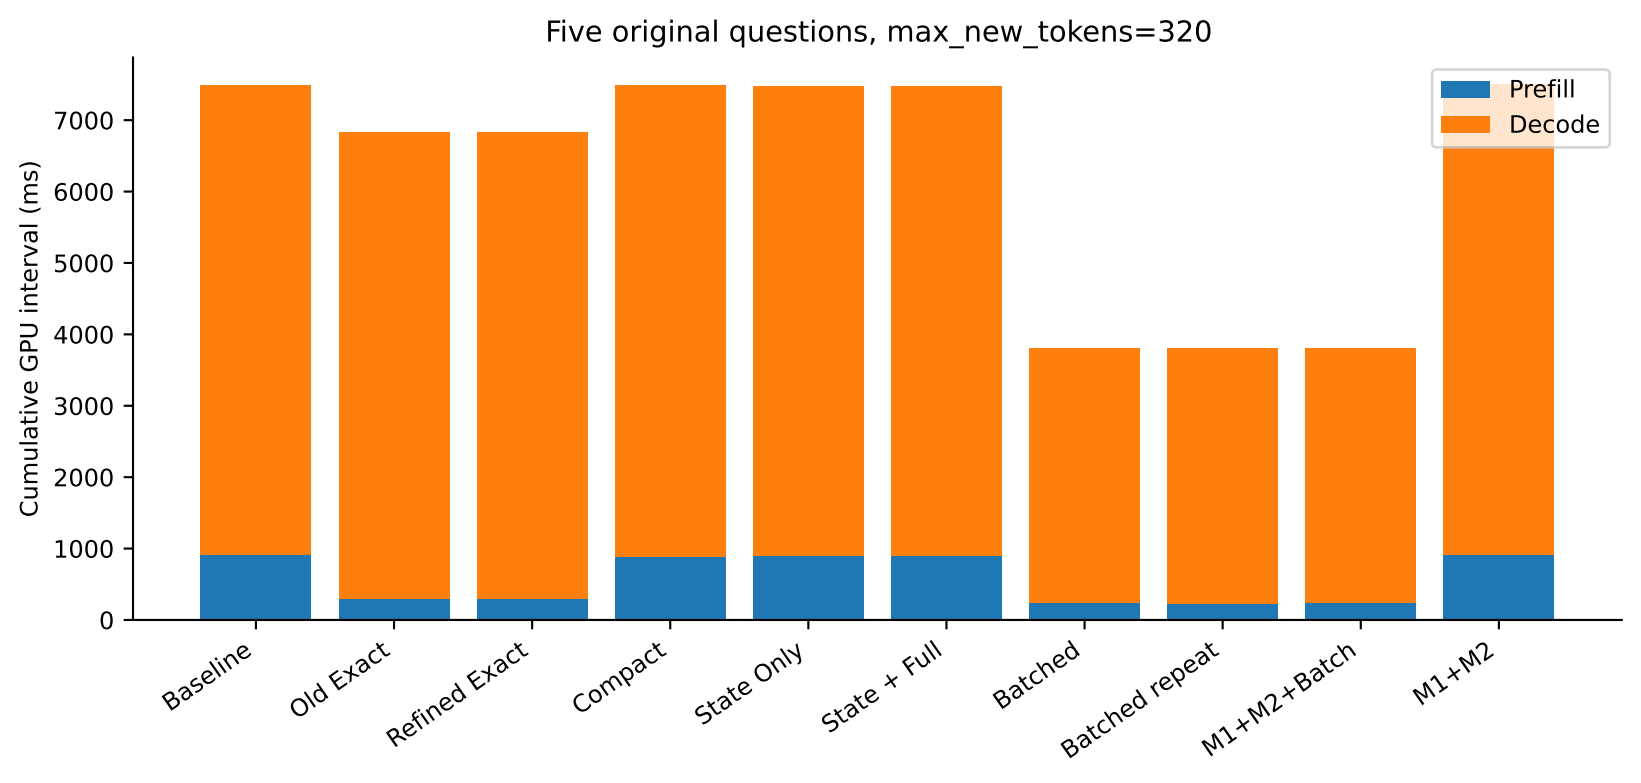
Five original questions, max_new_tokens=320
Cumulative GPU interval (ms)
0
1000
2000
3000
4000
5000
6000
7000
Prefill
Decode
Baseline
Old Exact
Refined Exact
Compact
State Only
State + Full
Batched
Batched repeat
M1+M2+Batch
M1+M2