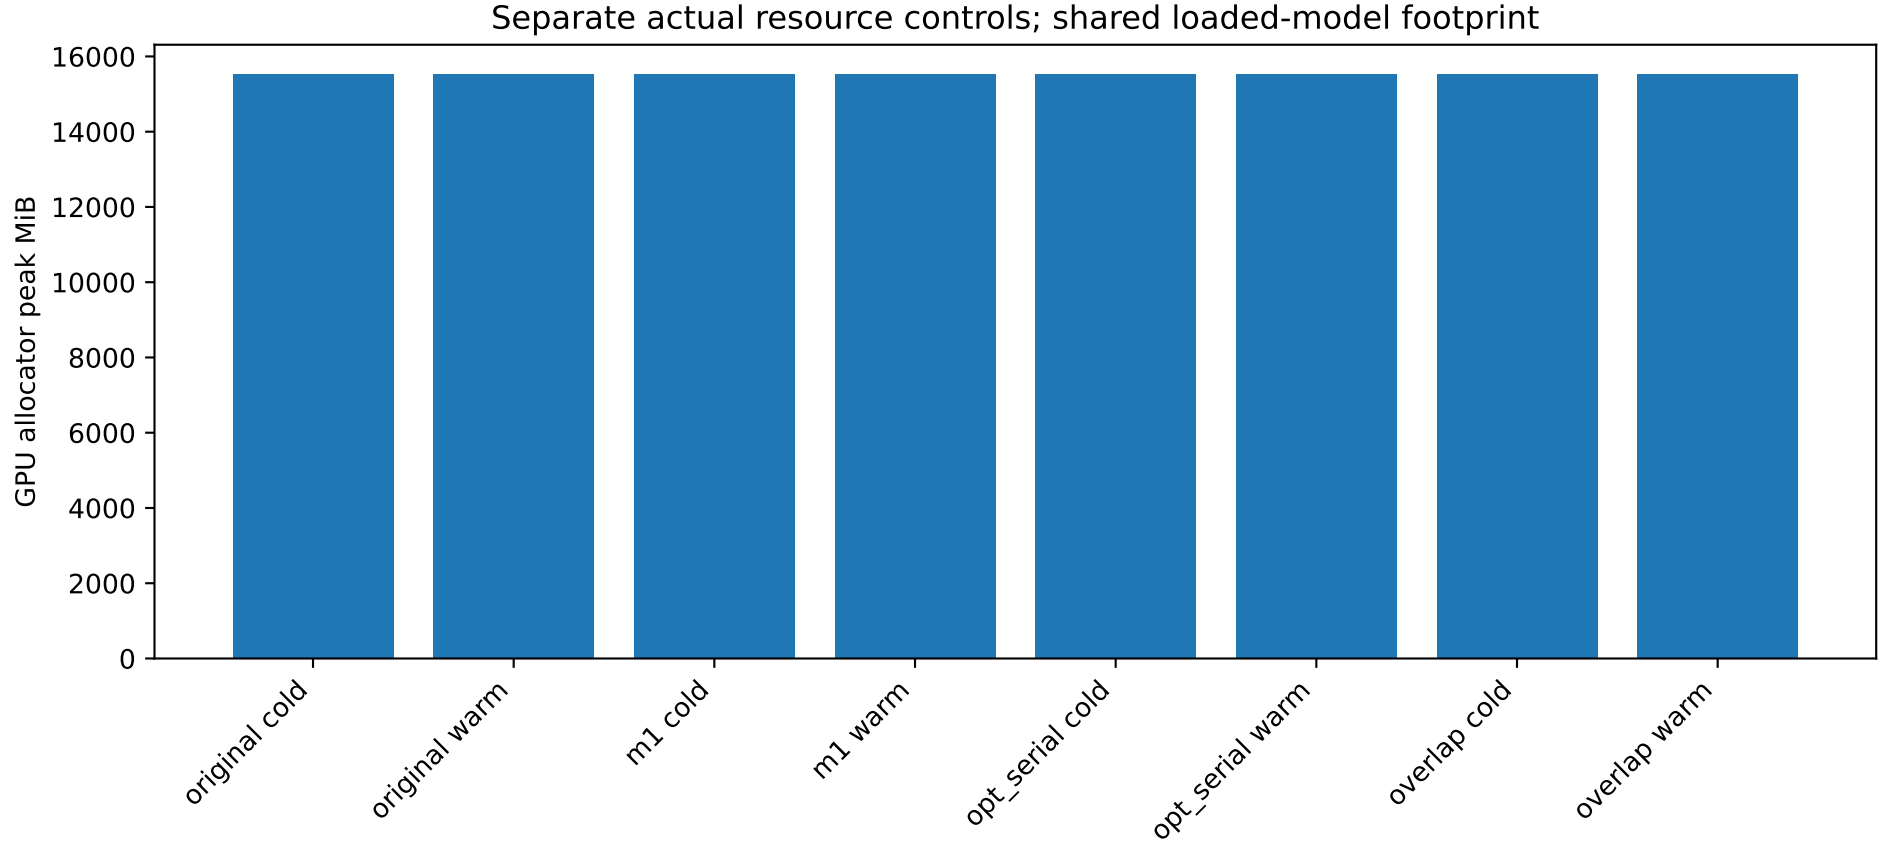
Separate actual resource controls; shared loaded-model footprint
GPU allocator peak MiB
0
2000
4000
6000
8000
10000
12000
14000
16000
original cold
original warm
m1 cold
m1 warm
opt_serial cold
opt_serial warm
overlap cold
overlap warm